

Diabetes_Prediction_Different_ML_Models

October 5, 2022

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

0.0.1 2.Loading the dataset

```
[2]: # loading the diabetes dataset to a pandas DataFrame
diabetes_dataset = pd.read_csv('diabetes.csv')
```

0.0.2 3. Exploratory Data Analysis

```
[3]: # printing the first 5 rows of the dataset
diabetes_dataset.head()
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[4]: # number of rows and Columns in this dataset
diabetes_dataset.shape
```

```
[4]: (768, 9)
```

```
[5]: #learning about the columns
diabetes_dataset.columns
```

```
[5]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
          'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
          dtype='object')
```

```
[6]: #knowledge of data type helps for computation
diabetes_dataset.dtypes
```

```
[6]: Pregnancies      int64
      Glucose          int64
      BloodPressure    int64
      SkinThickness    int64
      Insulin          int64
      BMI              float64
      DiabetesPedigreeFunction float64
      Age              int64
      Outcome          int64
      dtype: object
```

```
[7]: #Print a concise summary of a DataFrame
diabetes_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                    768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
[8]: # getting the statistical measures of the data
diabetes_dataset.describe()
```

```
[8]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	\
count	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	
std	3.369578	31.972618	19.355807	15.952218	115.244002	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	

75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

CONCLUSION :- We observe that min value of some columns is 0 which cannot be possible medically. Hence in the data cleaning process we'll have to replace them with median/mean value depending on the distribution. Also in the max column we can see insulin levels as high as 846! We have to treat outliers.

0.0.3 Data cleaning

- Dropping duplicate values
- Checking NULL values
- Checking for 0 value and replacing it :- It isn't medically possible for some data record to have 0 value such as Blood Pressure or Glucose levels. Hence we replace them with the mean value of that particular column.

```
[9]: #dropping duplicate values
diabetes_dataset = diabetes_dataset.drop_duplicates()
```

```
[10]: #check for missing values
diabetes_dataset.isnull().sum()
```

```
[10]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

```
[11]: #checking for 0 values in 5 columns
print(diabetes_dataset[diabetes_dataset['BloodPressure']==0].shape[0])
print(diabetes_dataset[diabetes_dataset['Glucose']==0].shape[0])
print(diabetes_dataset[diabetes_dataset['SkinThickness']==0].shape[0])
print(diabetes_dataset[diabetes_dataset['Insulin']==0].shape[0])
```

```
print(diabetes_dataset[diabetes_dataset['BMI']==0].shape[0])
```

```
35
5
227
374
11
```

```
[12]: #Age & DiabetesPedigreeFunction do not have have minimum 0 value so no need to_
      ↪replace
```

Some of the columns have a skewed distribution, so the mean is more affected by outliers than the median. Glucose and Blood Pressure have normal distributions hence we replace 0 values in those columns by mean value. SkinThickness, Insulin, BMI have skewed distributions hence median is a better choice as it is less affected by outliers.

```
[13]: #replacing 0 values with median of that column
diabetes_dataset['Glucose']=diabetes_dataset['Glucose'].
      ↪replace(0,diabetes_dataset['Glucose'].mean())
#normal distribution
diabetes_dataset['BloodPressure']=diabetes_dataset['BloodPressure'].
      ↪replace(0,diabetes_dataset['BloodPressure'].mean())
#normal distribution
diabetes_dataset['SkinThickness']=diabetes_dataset['SkinThickness'].
      ↪replace(0,diabetes_dataset['SkinThickness'].median())
#skewed distribution
diabetes_dataset['Insulin']=diabetes_dataset['Insulin'].
      ↪replace(0,diabetes_dataset['Insulin'].median())
#skewed distribution
diabetes_dataset['BMI']=diabetes_dataset['BMI'].
      ↪replace(0,diabetes_dataset['BMI'].median())
#skewed distribution
```

0.0.4 4.Data Visualization

Here we are going to plot :-


- Count Plot :- to see if the dataset is balanced or not
- Histograms :- to see if data is normally distributed or skewed
- Box Plot :- to analyse the distribution and see the outliers

```
[14]: diabetes_dataset['Outcome'].value_counts()
```

```
[14]: 0    500
      1    268
      Name: Outcome, dtype: int64
```

```
[15]: sns.countplot('Outcome',data=diabetes_dataset)
```

```
[15]: <AxesSubplot:xlabel='Outcome', ylabel='count'>
```

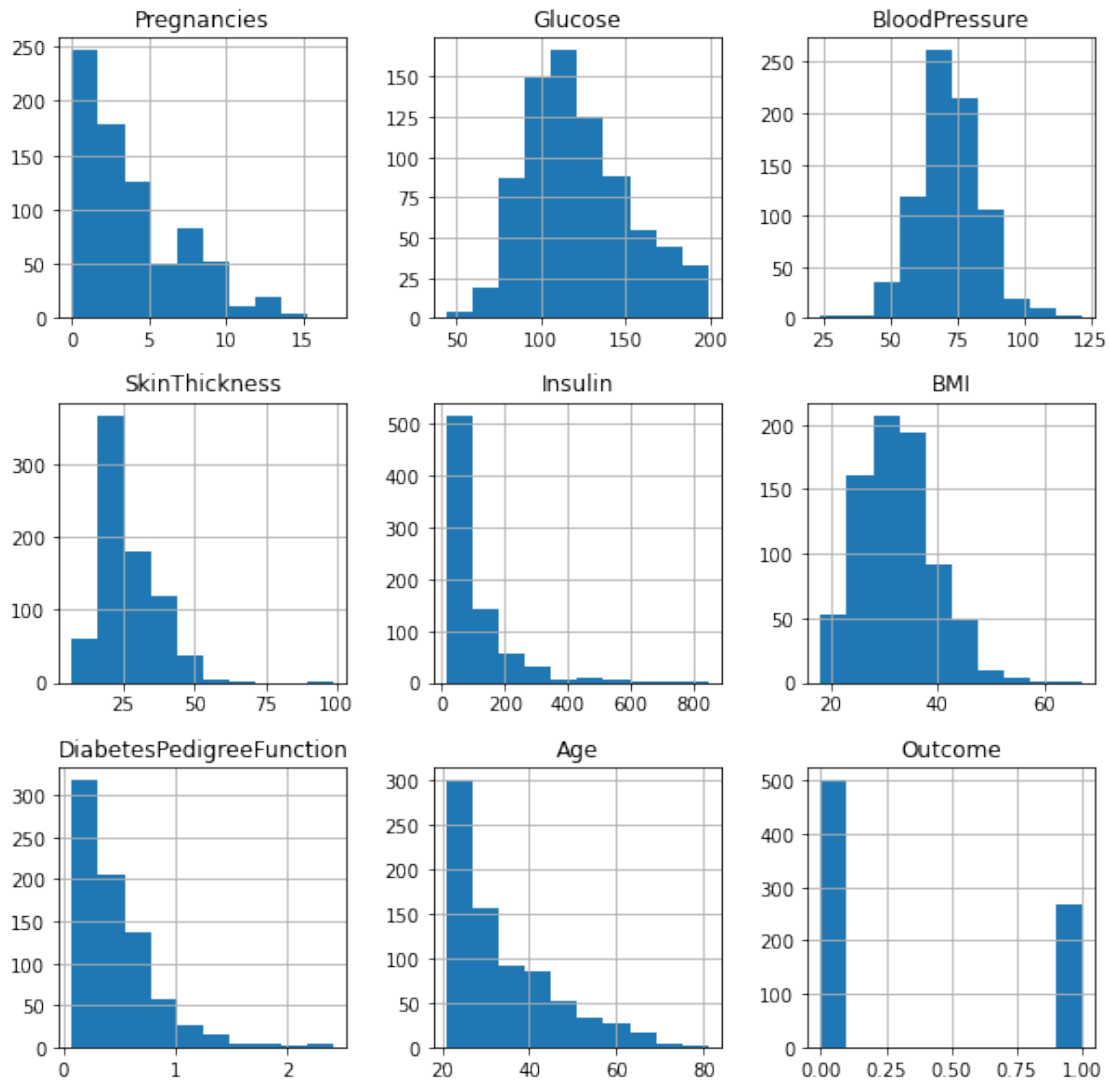


0 -> Non-Diabetic

1 -> Diabetic

Conclusion :- We observe that number of people who do not have diabetes is far more than people who do which indicates that our data is imbalanced.

```
[16]: #histogram for each feature  
diabetes_dataset.hist(bins=10,figsize=(10,10))  
plt.show()
```



Conclusion :- We observe that only glucose and Blood Pressure are normally distributed rest others are skewed and have outliers

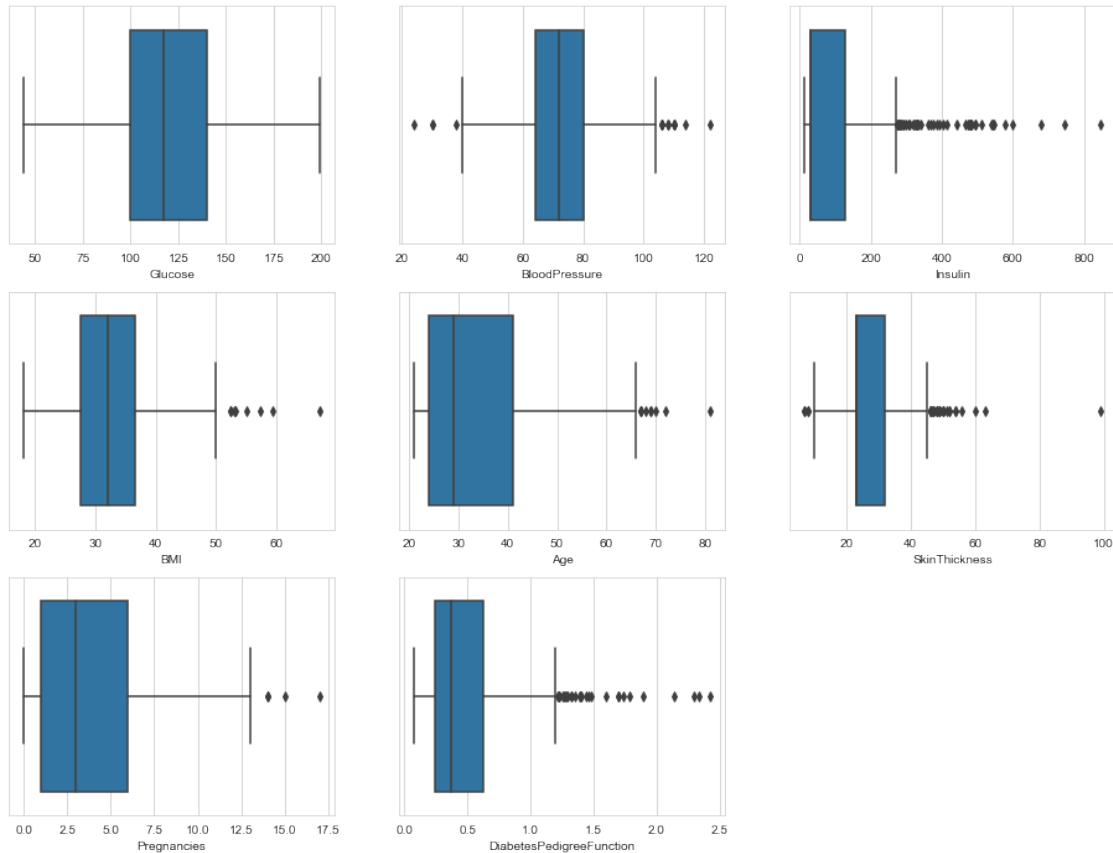
```
[17]: plt.figure(figsize=(16,12))
sns.set_style(style='whitegrid')
plt.subplot(3,3,1)
sns.boxplot(x='Glucose',data=diabetes_dataset)
plt.subplot(3,3,2)
sns.boxplot(x='BloodPressure',data=diabetes_dataset)
plt.subplot(3,3,3)
sns.boxplot(x='Insulin',data=diabetes_dataset)
plt.subplot(3,3,4)
sns.boxplot(x='BMI',data=diabetes_dataset)
plt.subplot(3,3,5)
```

```

sns.boxplot(x='Age',data=diabetes_dataset)
plt.subplot(3,3,6)
sns.boxplot(x='SkinThickness',data=diabetes_dataset)
plt.subplot(3,3,7)
sns.boxplot(x='Pregnancies',data=diabetes_dataset)
plt.subplot(3,3,8)
sns.boxplot(x='DiabetesPedigreeFunction',data=diabetes_dataset)

```

[17]: <AxesSubplot:xlabel='DiabetesPedigreeFunction'>




Outliers are unusual values in your dataset, and they can distort statistical analyses and violate their assumptions. Hence it is of most importance to deal with them. In this case removing outliers can cause data loss so we have to deal with it using various scaling and transformation techniques.

0.0.5 5.Feature Selection

```
[18]: corrmat=diabetes_dataset.corr()  
sns.heatmap(corrmat, annot=True)
```

[18]: <AxesSubplot:>



CONCLUSION :- Observe the last row 'Outcome' and note its correlation scores with different features. We can observe that Glucose, BMI and Age are the most correlated with Outcome. BloodPressure, Insulin, DiabetesPedigreeFunction are the least correlated, hence they don't contribute much to the model so we can drop them.

```
[19]: diabetes_selected=diabetes_dataset.  
      ↪drop(['BloodPressure', 'Insulin', 'DiabetesPedigreeFunction'],axis='columns')
```

0.0.6 6.handling Outliers

```
[20]: from sklearn.preprocessing import QuantileTransformer  
x=diabetes_selected  
quantile = QuantileTransformer()
```



```

X = quantile.fit_transform(x)
diabetes_new=quantile.transform(X)
diabetes_new=pd.DataFrame(X)
diabetes_new.columns =['Pregnancies',
↳ 'Glucose', 'SkinThickness', 'BMI', 'Age', 'Outcome']
diabetes_new.head(10)

```

```

[20]:
   Pregnancies  Glucose  SkinThickness    BMI    Age  Outcome
0      0.747718  0.810300      0.801825  0.591265  0.889831      1.0
1      0.232725  0.091265      0.644720  0.213168  0.558670      0.0
2      0.863755  0.956975      0.357888  0.077575  0.585398      1.0
3      0.232725  0.124511      0.357888  0.284224  0.000000      0.0
4      0.000000  0.721643      0.801825  0.926988  0.606258      1.0
5      0.677966  0.483703      0.357888  0.171447  0.529335      0.0
6      0.503259  0.039765      0.735332  0.433507  0.368970      1.0
7      0.940678  0.473272      0.357888  0.691656  0.496741      0.0
8      0.387223  0.996089      0.956323  0.408083  0.919166      1.0
9      0.863755  0.604302      0.357888  0.487614  0.926336      1.0

```

```


[21]: plt.figure(figsize=(16,12))
sns.set_style(style='whitegrid')
plt.subplot(3,3,1)
sns.boxplot(x=diabetes_new['Glucose'],data=diabetes_new)
plt.subplot(3,3,2)
sns.boxplot(x=diabetes_new['BMI'],data=diabetes_new)
plt.subplot(3,3,3)
sns.boxplot(x=diabetes_new['Pregnancies'],data=diabetes_new)
plt.subplot(3,3,4)
sns.boxplot(x=diabetes_new['Age'],data=diabetes_new)
plt.subplot(3,3,5)
sns.boxplot(x=diabetes_new['SkinThickness'],data=diabetes_new)

```

```

[21]: <AxesSubplot:xlabel='SkinThickness'>

```



0.0.7 7.Split the Data Frame into X and Y

```
[22]: target_name='Outcome'
y= diabetes_new[target_name]#given predictions - training data
X=diabetes_new.drop(target_name,axis=1)
```

```
[23]: X.head()
```

```
[23]:
```

	Pregnancies	Glucose	SkinThickness	BMI	Age
0	0.747718	0.810300	0.801825	0.591265	0.889831
1	0.232725	0.091265	0.644720	0.213168	0.558670
2	0.863755	0.956975	0.357888	0.077575	0.585398
3	0.232725	0.124511	0.357888	0.284224	0.000000
4	0.000000	0.721643	0.801825	0.926988	0.606258

```
[24]: y.head()
```

```
[24]:
```

0	1.0
1	0.0
2	1.0
3	0.0
4	1.0

Name: Outcome, dtype: float64

0.0.8 8.Train Test Split

```
[25]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test= train_test_split(X,y,test_size=0.
↪2,random_state=0)
#splitting data in 80% train, 20%test
```

```
[26]: X_train.shape,y_train.shape
```

```
[26]: ((614, 5), (614,))
```

```
[27]: X_test.shape,y_test.shape
```

```
[27]: ((154, 5), (154,))
```

0.0.9 9.Classification Algorithms

- **9.1 K Nearest Neighbours** :- KNN algorithm, is a non-parametric algorithm that classifies data points based on their proximity and association to other available data.

```
[28]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import RepeatedStratifiedKFold
from sklearn.metrics import classification_report,confusion_matrix
from sklearn.metrics import f1_score, precision_score, recall_score
from sklearn.model_selection import GridSearchCV
```

```
[29]: #List Hyperparameters to tune
knn= KNeighborsClassifier()
n_neighbors = list(range(15,25))
p=[1,2]
weights = ['uniform', 'distance']
metric = ['euclidean', 'manhattan', 'minkowski']

#convert to dictionary
hyperparameters = dict(n_neighbors=n_neighbors,
↪p=p,weights=weights,metric=metric)

#Making model
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
grid_search = GridSearchCV(estimator=knn, param_grid=hyperparameters,
↪n_jobs=-1, cv=cv, scoring='f1',error_score=0)
```

```
[30]: best_model = grid_search.fit(X_train,y_train)
```

```
[31]: #Best Hyperparameters Value
print('Best leaf_size:', best_model.best_estimator_.get_params()['leaf_size'])
print('Best p:', best_model.best_estimator_.get_params()['p'])
```

```
print('Best n_neighbors:', best_model.best_estimator_.
      ↪get_params()['n_neighbors'])
```

Best leaf_size: 30
 Best p: 1
 Best n_neighbors: 19

```
[32]: #Predict testing set
knn_pred = best_model.predict(X_test)
```

```
[33]: print("\n Confusion Matrix:\n",confusion_matrix(y_test,knn_pred))
print("\n Classification Report is:\n",classification_report(y_test,knn_pred))
print("\n F1:\n",f1_score(y_test,knn_pred))
print("\n Precision score is:\n",precision_score(y_test,knn_pred))
print("\n Recall score is:\n",recall_score(y_test,knn_pred))
```

Confusion Matrix:

```
[[94 13]
 [17 30]]
```

Classification Report is:

	precision	recall	f1-score	support
0.0	0.85	0.88	0.86	107
1.0	0.70	0.64	0.67	47
accuracy			0.81	154
macro avg	0.77	0.76	0.76	154
weighted avg	0.80	0.81	0.80	154

F1:

0.6666666666666666

Precision score is:

0.6976744186046512

Recall score is:

0.6382978723404256

- **9.2 Naive Bayes** :- Naive Bayes is classification approach that adopts the principle of class conditional independence from the Bayes Theorem. This means that the presence of one feature does not impact the presence of another in the probability of a given outcome, and each predictor has an equal effect on that result

```
[34]: from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import GridSearchCV
```

```
param_grid_nb = {
    'var_smoothing': np.logspace(0,-2, num=100)
}
nbModel_grid = GridSearchCV(estimator=GaussianNB(), param_grid=param_grid_nb,
    ↪ verbose=1, cv=10, n_jobs=-1)
```

```
[35]: best_model= nbModel_grid.fit(X_train, y_train)
```

Fitting 10 folds for each of 100 candidates, totalling 1000 fits

```
[36]: nb_pred=best_model.predict(X_test)
```

```
[37]: print("\n Confusion Matrix:\n", confusion_matrix(y_test,nb_pred))
print("\n Classification Report is:\n",classification_report(y_test,nb_pred))
print("\n F1:\n",f1_score(y_test,nb_pred))
print("\n Precision score is:\n",precision_score(y_test,nb_pred))
print("\n Recall score is:\n",recall_score(y_test,nb_pred))
```

Confusion Matrix:

```
[[93 14]
 [22 25]]
```

Classification Report is:

	precision	recall	f1-score	support
0.0	0.81	0.87	0.84	107
1.0	0.64	0.53	0.58	47
accuracy			0.77	154
macro avg	0.72	0.70	0.71	154
weighted avg	0.76	0.77	0.76	154

F1:

```
0.5813953488372093
```

Precision score is:

```
0.6410256410256411
```

Recall score is:

```
0.5319148936170213
```

- **9.3 Support Vector Machine :-** It is typically leveraged for classification problems, constructing a hyperplane where the distance between two classes of data points is at its maximum. This hyperplane is known as the decision boundary, separating the classes of data points (e.g., has diabetes vs doesn't have diabetes) on either side of the plane.

```
[38]: from sklearn.model_selection import RepeatedStratifiedKFold
      from sklearn.model_selection import GridSearchCV
      from sklearn.svm import SVC
      from sklearn.metrics import classification_report, confusion_matrix
      from sklearn.metrics import f1_score, precision_score, recall_score

[39]: model = SVC()
      kernel = ['poly', 'rbf', 'sigmoid']
      C = [50, 10, 1.0, 0.1, 0.01]
      gamma = ['scale']

[40]: # define grid search
      grid = dict(kernel=kernel, C=C, gamma=gamma)
      cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
      grid_search = GridSearchCV(estimator=model, param_grid=grid, n_jobs=-1, cv=cv,
      ↪scoring='f1', error_score=0)

[41]: grid_result = grid_search.fit(X, y)

[42]: svm_pred=grid_result.predict(X_test)

[43]: print("\n Confusion Matrix:\n", confusion_matrix(y_test,svm_pred))
      print("\n Classification Report is:\n",classification_report(y_test,svm_pred))
      print("\n F1:\n",f1_score(y_test,knn_pred))
      print("\n Precision score is:\n",precision_score(y_test,knn_pred))
      print("\n Recall score is:\n",recall_score(y_test,knn_pred))
```

Confusion Matrix:

```
[[95 12]
 [15 32]]
```

Classification Report is:

	precision	recall	f1-score	support
0.0	0.86	0.89	0.88	107
1.0	0.73	0.68	0.70	47
accuracy			0.82	154
macro avg	0.80	0.78	0.79	154
weighted avg	0.82	0.82	0.82	154

F1:

0.6666666666666666

Precision score is:

0.6976744186046512

Recall score is:
0.6382978723404256

- **9.4 Logistic Regression** :- Logistical regression is selected when the dependent variable is categorical, meaning they have binary outputs, such as “true” and “false” or “yes” and “no.”
- Logistic regression does not really have any critical hyperparameters to tune. Sometimes, you can see useful differences in performance or convergence with different solvers (solver). Regularization (penalty) can sometimes be helpful.

```
[44]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import classification_report, confusion_matrix
      from sklearn.metrics import f1_score, precision_score, recall_score, accuracy_score
```

```
[45]: reg = LogisticRegression()
      reg.fit(X_train, y_train)
```

```
[45]: LogisticRegression()
```

```
[46]: lr_pred = reg.predict(X_test)
```

```
[47]: print("\n Confusion Matrix:\n", confusion_matrix(y_test, lr_pred))
      print("\n Classification Report is:\n", classification_report(y_test, lr_pred))
      print("\n F1:\n", f1_score(y_test, lr_pred))
      print("\n Precision score is:\n", precision_score(y_test, lr_pred))
      print("\n Recall score is:\n", recall_score(y_test, lr_pred))
```

Confusion Matrix:

```
[[95 12]
 [20 27]]
```

Classification Report is:

	precision	recall	f1-score	support
0.0	0.83	0.89	0.86	107
1.0	0.69	0.57	0.63	47
accuracy			0.79	154
macro avg	0.76	0.73	0.74	154
weighted avg	0.79	0.79	0.79	154

F1:

0.627906976744186

Precision score is:

0.6923076923076923

Recall score is:
0.574468085106383

0.0.10 Accuracy Score

K Nearest Neighbours :- 81%

Naive Bayes :- 77%

Support Vector Machine :- 82%

Logistic Regression :- 79%